

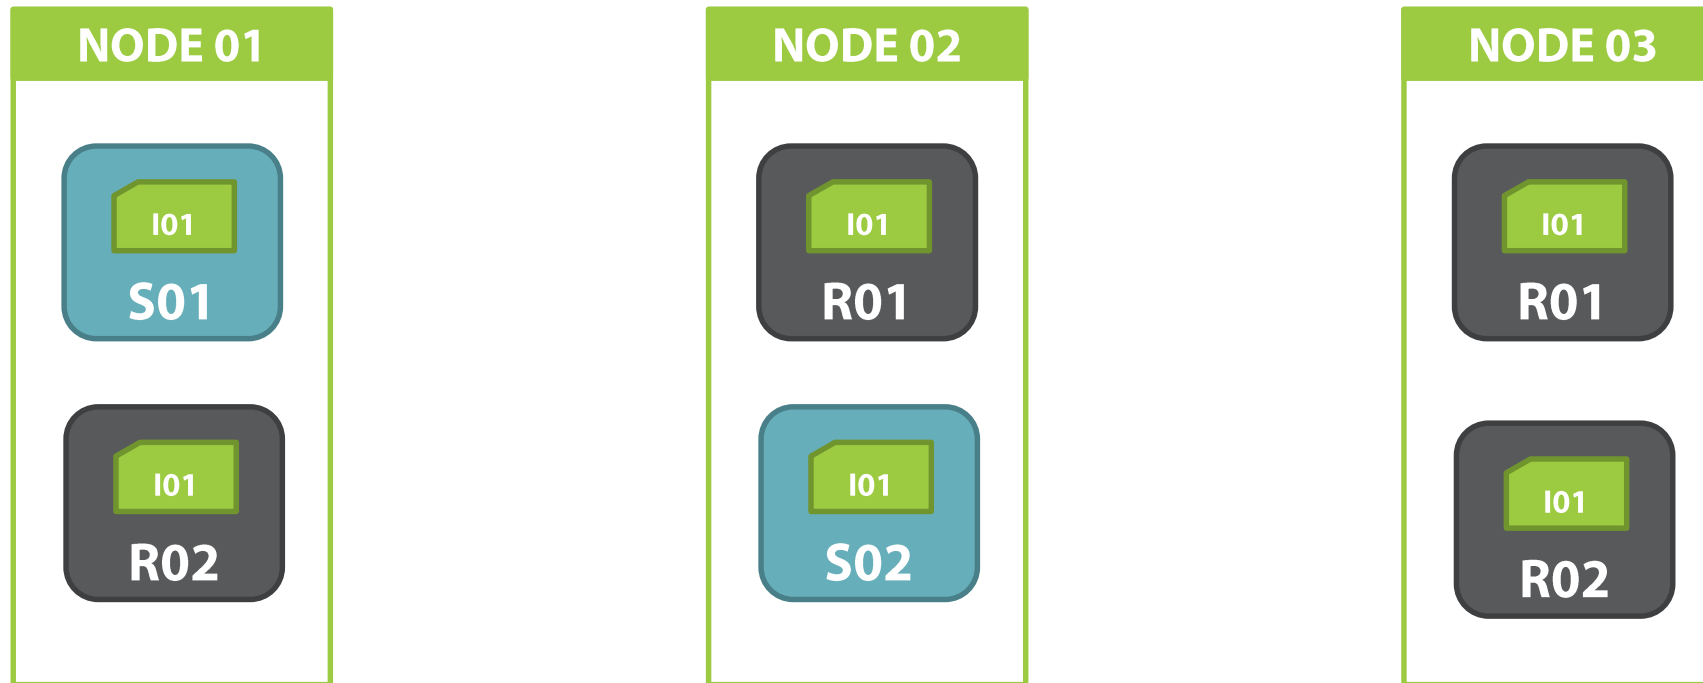
Index Strategies



JP Toto

@jptoto | <http://jptoto.jp>

Indexes are Stored in Shards



How do we actually create indexes?

There's an http API for that!

Creating an Index

POST http://es-client-01:9200/knowledge_base

```
{
  "settings": {
    "index": {
      "number_of_shards": 5,
      "number_of_replicas": 1
    }
  },
  "mappings": {
    "kb_article": {
      "properties": {
        "article_name": {
          "type": "string"
        },
        "article_text": {
          "type": "string"
        },
        "article_date": {
          "type": "date"
        }
      }
    }
  }
}
```

Shard Size Considerations



Talk to data owners to gauge document size

Perform load tests on sample data

Large shards take longer to recreate

Rolling Indexes



globo_index

Static index


log_20151125

log_20151126

log_20151127

Daily "rolling" indexes

Rolling Indexes: Logging

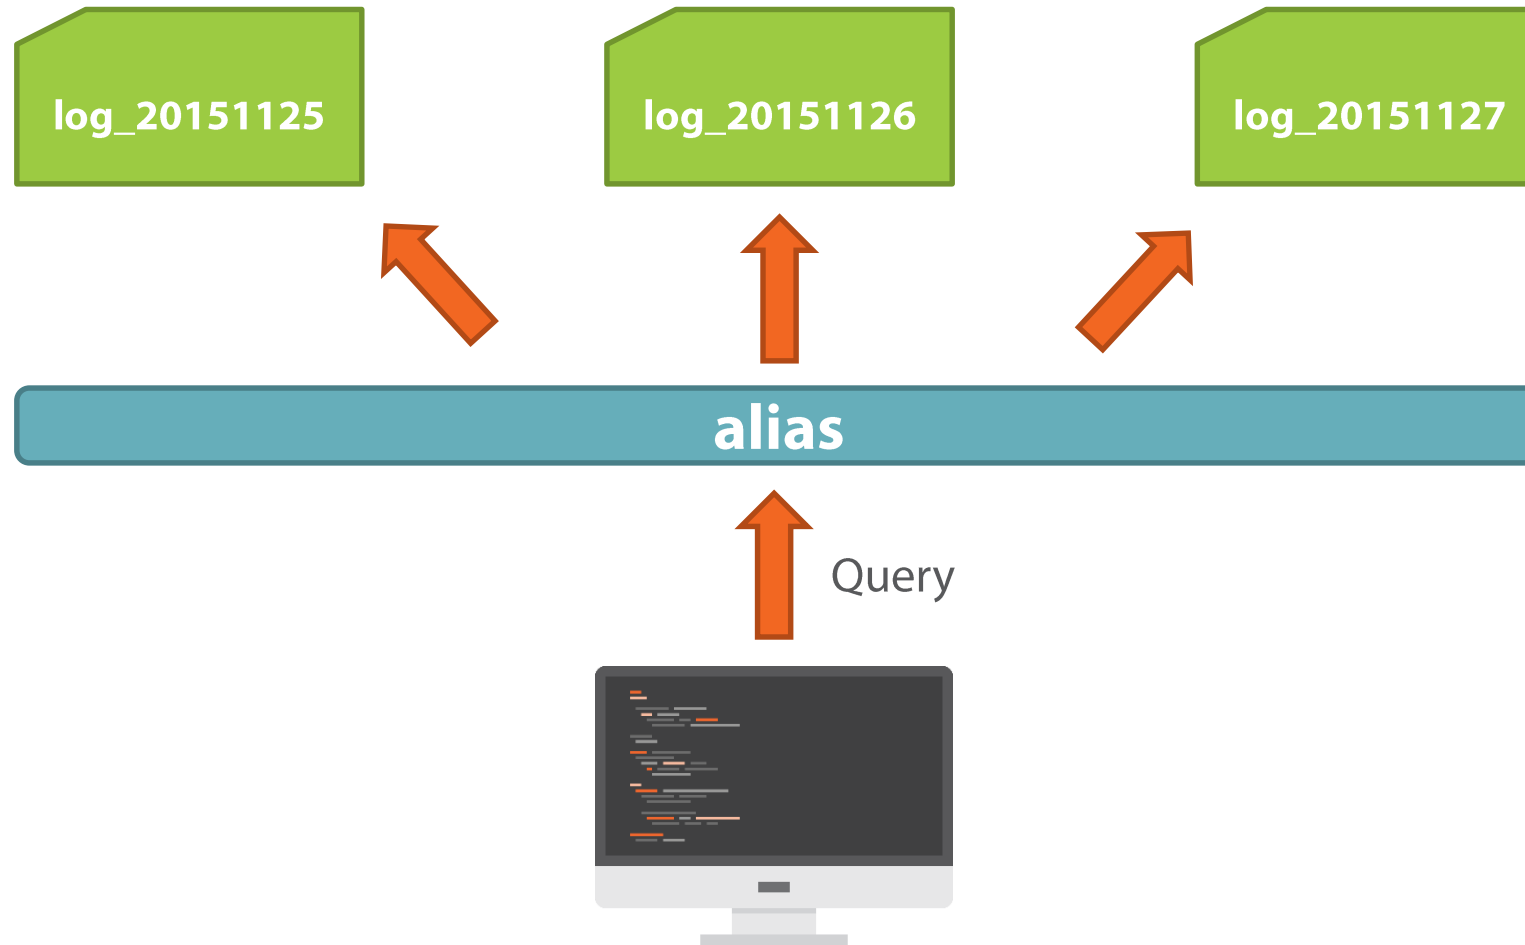


Introducing : Templates and Aliases

```
{
  "template": "log_*",
  "aliases": {
    "log": {}
  },
  "settings": {
    "index": {
      "number_of_shards": 5,
      "number_of_replicas": 1
    }
  },
  "mappings": ...
}
```

Template

Introducing : Templates and Aliases



Introducing : Templates and Aliases

```
{
  "template": "log_*",
  "aliases": {
    "log": {}
  },
  "settings": {
    "index": {
      "number_of_shards": 5,
      "number_of_replicas": 1
    }
  },
  "mappings": {
    "log_event": {
      "properties": {
        "log_name": {
          "type": "string"
        },
        "log_text": {
          "type": "string"
        },
        "log_date": {
          "type": "date"
        }
      }
    }
  }
}
```

alias

log_20151125

log_20151126

log_20151127

Dynamic Indexes



log_20151125

log_20151126

string ?

int ?

date ?

log_20151127

TEMPLATE



Indexes and Shards



Learned how indexes are created

Explored time-based “rolling” indexes

Disk considerations for shard allocation